

HW 1

Bryan Mui - UID 506021334 - 14 April 2025

Problem 1

Part a

Find the theoretical min for the function:

$$f(x) = x^4 + 2x^2 + 1$$

Solution: find $f'(x)$ and $f''(x)$, set $f'(x)$ to 0 and solve, and $f''(x)$ needs to be > 0 to be a min

Step 1: find $f'(x)$ and $f''(x)$

$$f(x) = x^4 + 2x^2 + 1 \quad (1)$$

$$f'(x) = 4x^3 + 4x \quad (2)$$

$$f''(x) = 12x^2 + 4 \quad (3)$$

$$(4)$$

Step 2: set $f'(x)$ to 0 and solve

$$f'(x) = 4x^3 + 4x \quad (5)$$

$$0 = 4x^3 + 4x \quad (6)$$

$$0 = 4x(x^2 + 4) \quad (7)$$

We get

$$x = 0$$

and

$$0 = x^2 + 4$$

which has no real solution

Step 3: check that $f''(x)$ needs to be > 0 to be a min

Our critical point is $x = 0$,

$$f''(0) = 12(0)^2 + 4 \quad (8)$$

$$= 4 \quad (9)$$

Since $f'(x) = 0$ at 0 and $f''(x) > 0$ at that point, **we have a min at $x = 0$, and plugging into $f(0)$ we get the minimum point**

$$(0, 1)$$

Part b

0)

Use the gradient descent algorithm with **constant step size** and with **back-tracking line search** to calculate x_{min}

Constant step size descent is implemented as follows:

1. Select a random starting point x_0
2. While stopping criteria $<$ tolerance, do:
 - Select η_k (as a constant)
 - Calculate $x_{(k+1)} = x_k - \eta_k * \nabla(f(x_k))$
 - Calculate the value of stopping criterion

Stopping criteria: Stop if $\|\nabla(f(x_k))\|_2 \leq \epsilon$

```
# Gradient descent algorithm that uses backtracking to minimize an objective
↪ function

gradient_descent_constant_step <- function(tol = 1e-6, max_iter = 10000,
↪ step_size = 0.01) {
  # Step 1: Initialize and select a random stopping point
  # Initialize
```

```

set.seed(777) # example seeding
last_iter <- 0 # the last iteration ran
eta <- step_size # step size that is decided manually
max_iter <- max_iter # max iterations before terminating if minimum isn't
↪ found
tolerance <- tol # tolerance for the stopping criteria
obj_values <- numeric(max_iter) # Stores the value of f(x)
eta_values <- numeric(max_iter) # To store eta values used each iteration
betas <- numeric(max_iter) # Stores the value of x guesses
x0 <- runif(1, min=-10, max=10) # our first guess is somewhere between
↪ -10-10

# Set the objective function to the function to be minimized
# Objective function: f(x)
obj_function <- function(x) {
  return(x^4 + 2*(x^2) + 1)
}

# Gradient function: d/dx of f(x)
gradient <- function(x) {
  return(4*x^3 + 4*x)
}

# Append the first guess to the obj_values and betas vector
betas[1] <- x0
obj_values[1] <- obj_function(x0)

# Step 2: While stopping criteria < tolerance, do:
for (iter in 1:max_iter) { # the iteration goes n = 1, 2, 3, 4, but the
↪ arrays of our output starts at iter = 0 and guess x0
  # Select eta(step size), which is constant
  # There's nothing to do for this step

  # Calculate the next guess of x_k+1, calculate f(x_k+1), set eta(x_k+1)
  betas[iter + 1] <- betas[iter] - (eta * gradient(betas[iter]))
  obj_values[iter + 1] <- obj_function(betas[iter + 1])
  eta_values[iter + 1] <- eta

  # Calculate the value of the stopping criterion
  stop_criteria <- abs(gradient(betas[iter + 1]))

  # If stopping criteria less than tolerance, break
  if(is.na(stop_criteria) || stop_criteria <= tolerance) {

```

```

    last_iter <- iter + 1
    break
  }

  # if we never set last iter, then we hit the max number of iterations and
  ↪ need to set
  if(last_iter == 0) { last_iter <- max_iter }

  # end algorithm
}

return(list(betas = betas, obj_values = obj_values, eta_values =
  ↪ eta_values, last_iter = last_iter)) # in this case, beta(predictors)
  ↪ are the x values, obj_values are f(x), eta is the step size, last iter
  ↪ is the value in the vector of the final iteration before stopping
}

```

Running the gradient descent algorithm with fixed step size:

```

minimize <- gradient_descent_constant_step(tol = 1e-6, max_iter = 10000,
  ↪ step_size = 0.01)
print(minimize)

```

\$betas

```

[1] 3.7571481 1.4854017 1.2948890 1.1562458 1.0481644 0.9601753 0.8863594
[8] 0.8230509 0.7678271 0.7190068 0.6753783 0.6360406 0.6003067 0.5676411
[15] 0.5376194 0.5098990 0.4842001 0.4602913 0.4379788 0.4170990 0.3975125
[22] 0.3790995 0.3617562 0.3453923 0.3299284 0.3152947 0.3014292 0.2882765
[29] 0.2757872 0.2639167 0.2526247 0.2418748 0.2316338 0.2218713 0.2125596
[36] 0.2036731 0.1951882 0.1870832 0.1793380 0.1719337 0.1648531 0.1580797
[43] 0.1515985 0.1453952 0.1394565 0.1337697 0.1283232 0.1231057 0.1181069
[50] 0.1133167
[ reached 'max' / getOption("max.print") -- omitted 9950 entries ]

```

\$obj_values

```

[1] 228.498357 10.281118 7.164923 5.461122 4.404326 3.693840
[7] 3.188485 2.813714 2.526696 2.301200 2.120332 1.972754
[13] 1.850601 1.748256 1.661610 1.587592 1.523866 1.468624
[19] 1.420448 1.378209 1.341002 1.308087 1.278861 1.252823
[25] 1.229554 1.208704 1.189975 1.173113 1.157902 1.144155
[31] 1.131711 1.120430 1.110187 1.100877 1.092405 1.084686

```

```

[37] 1.077648 1.071225 1.065359 1.059996 1.055092 1.050603
[43] 1.046492 1.042726 1.039274 1.036109 1.033205 1.030540
[49] 1.028093 1.025846
[ reached 'max' / getOption("max.print") -- omitted 9950 entries ]

$eta_values
[1] 0.00 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
[16] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
[31] 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01 0.01
[46] 0.01 0.01 0.01 0.01 0.01
[ reached 'max' / getOption("max.print") -- omitted 9950 entries ]

$last_iter
[1] 369

```

```
cat("The functions stopped after", minimize$last_iter - 1, "iterations \n")
```

The functions stopped after 368 iterations

```

cat("The function's point of minimization is", "(",
  ↪ minimize$betas[minimize$last_iter], ",",
  ↪ minimize$obj_values[minimize$last_iter], ") \n")

```

The function's point of minimization is (2.488126e-07 , 1)

Backtracking Line Search is implemented as follows:

1. Select a random starting point x_0
2. While stopping criteria $<$ tolerance, do:
 - Select η_k using backtracking line search
 - Calculate $x_{(k+1)} = x_k - \eta_k * \nabla(f(x_k))$
 - Calculate the value of stopping criterion

Backtracking Line Search:

- Set $\eta^0 > 0$ (usually a large value), $\epsilon \in (0, 1)$ and $\tau \in (0, 1)$
- Set $\eta_1 = \eta^0$
- At iteration k , set $\eta_k < -\eta_{k-1}$

1. Check whether the Armijo Condition holds:

$$h(\eta_k) \leq h(0) + \epsilon \eta_k h'(0)$$

where $h(\eta_k) = f(x_k) - \eta_k \nabla f(x_k)$

2.

- If yes(condition holds), terminate and keep η_k
- If no, set $\eta_k = \tau \eta_k$ and go to Step 1

Stopping criteria: Stop if $\|\nabla f(x_k)\|_2 \leq \epsilon$

Other note: Since we need $h'(0)$ for the Armijo condition calculation, that is given by:

$$h'(0) = -[\nabla f(x_k)]^\top \nabla f(x_k)$$

Since we are minimizing x , a one dimensional function, we can simplify to

$$h'(0) = -\|\nabla f(x_k)\|^2$$

To summarize, backtracking line search chooses the step size by ensuring the Armijo condition always holds. If the Armijo condition doesn't hold, we are probably overshooting, hence the step size gets updated iteratively

```
gradient_descent_backtracking <- function(tol = 1e-6, max_iter = 10000,
  ↪  epsilon = 0.5, tau = 0.8) {
  # Step 1: Initialize and select a random stopping point
  # Initialize
  set.seed(777) # example seeding
  last_iter <- 0 # the last iteration ran
  max_iter <- max_iter # max iterations before terminating if minimum isn't
  ↪  found
  tolerance <- tol # tolerance for the stopping criteria
  epsilon <- epsilon # Epsilon used in the step size criteria calculation
  tau <- tau # tau used in the step size criteria calculation
  obj_values <- numeric(max_iter) # Stores the value of f(x)
  eta_values <- numeric(max_iter) # To store eta values used each iteration
  betas <- numeric(max_iter) # Stores the value of x guesses
  x0 <- runif(1, min=-10, max=10) # our first guess is somewhere between -10
  ↪  to 10
  eta <- 1 # our initial step size

  # Set the objective function to the function to be minimized
  # Objective function: f(x)
```

```

obj_function <- function(x) {
  return(x^4 + 2*(x^2) + 1)
}

# Gradient function: d/dx of f(x)
gradient <- function(x) {
  return(4*x^3 + 4*x)
}

# Armijo condition function
armijo <- function(eta_k) {

}

# Append the first guess to the obj_values and betas vector
betas[1] <- x0
obj_values[1] <- obj_function(x0)

# Step 2: While stopping criteria < tolerance, do:
for (iter in 1:max_iter) { # the iteration goes n = 1, 2, 3, 4, but the
  ↪ arrays of our output starts at iter = 0 and guess x0
  # Select eta(step size) using backtracking line search
  if(iter > 1) {

  }

  # Calculate the next guess of x_k+1, calculate f(x_k+1), set eta(x_k+1)
  betas[iter + 1] <- betas[iter] - (eta * gradient(betas[iter]))
  obj_values[iter + 1] <- obj_function(betas[iter + 1])
  eta_values[iter + 1] <- eta

  # Calculate the value of the stopping criterion
  stop_criteria <- abs(gradient(betas[iter + 1]))

  # If stopping criteria less than tolerance, break
  if(is.na(stop_criteria) || stop_criteria <= tolerance) {
    last_iter <- iter + 1
    break
  }

  # if we never set last iter, then we hit the max number of iterations and
  ↪ need to set
  if(last_iter == 0) { last_iter <- max_iter }
}

```

```

    # end algorithm
}

return(list(betas = betas, obj_values = obj_values, eta_values =
  ↪ eta_values, last_iter = last_iter)) # in this case, beta(predictors)
  ↪ are the x values, obj_values are f(x), eta is the step size, last iter
  ↪ is the value in the vector of the final iteration before stopping
}

```

1) For the constant step size version of gradient descent, discuss how you selected the step size used in your code

Theoretical Analysis proves that for functions with a unique global minimum, the step size should be within 0 to $1/L$ to converge to the unique global minimum

2) For both versions of the gradient descent algorithm, plot the value of $f(x_k)$ as a function of k the number of iterations

3) or the the gradient descent method with backtracking line search, plot the step size η_k selected at step k as a function of k. Comment on the result